

# Investigate\_a\_Dataset

October 19, 2018

## 1 Project: Investigate a Movies Dataset

### 1.1 Table of Contents

Introduction

    Data Wrangling

    Exploratory Data Analysis

    Conclusions

## Introduction Investigate a movie dataset to find the common thing for a good movie as well as some of their characteristic to be able to make the highest rating or profit. Here are my curious question that I am seeking to answer based on this data 1. How did the customer rating of those movies change from the past to present? 2. What is the sweet spot for a movie length? (of which they make the most revenue)

    Firsly, I imported some libraries that needed for the analysis

```
In [19]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
% matplotlib inline
```

## Data Wrangling

#### 1.1.1 General Properties

Secondly, we read the data from CSV file, then we take a look at some of the rows in the top of the files

```
In [20]: df = pd.read_csv('tmdb-movies.csv')
df.head()
```

```
Out[20]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title \			
0	Jurassic World			
1	Mad Max: Fury Road			
2	Insurgent			
3	Star Wars: The Force Awakens			
4	Furious 7			
		cast \		
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...			
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...			
2	Shailene Woodley Theo James Kate Winslet Ansel...			
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...			
4	Vin Diesel Paul Walker Jason Statham Michelle ...			
		homepage	director \	
0	<a href="http://www.jurassicworld.com/">http://www.jurassicworld.com/</a>	Colin Trevorrow		
1	<a href="http://www.madmaxmovie.com/">http://www.madmaxmovie.com/</a>	George Miller		
2	<a href="http://www.thedivergentseries.movie/#insurgent">http://www.thedivergentseries.movie/#insurgent</a>	Robert Schwentke		
3	<a href="http://www.starwars.com/films/star-wars-episod...">http://www.starwars.com/films/star-wars-episod...</a>	J.J. Abrams		
4	<a href="http://www.furious7.com/">http://www.furious7.com/</a>	James Wan		
		tagline	...	\
0	The park is open.	...		
1	What a Lovely Day.	...		
2	One Choice Can Destroy You	...		
3	Every generation has a story.	...		
4	Vengeance Hits Home	...		
		overview runtime	\	
0	Twenty-two years after the events of Jurassic ...	124		
1	An apocalyptic story set in the furthest reach...	120		
2	Beatrice Prior must confront her inner demons ...	119		
3	Thirty years after defeating the Galactic Empi...	136		
4	Deckard Shaw seeks revenge against Dominic Tor...	137		
		genres \		
0	Action Adventure Science Fiction Thriller			
1	Action Adventure Science Fiction Thriller			
2	Adventure Science Fiction Thriller			
3	Action Adventure Science Fiction Fantasy			
4	Action Crime Thriller			
		production_companies	release_date	vote_count \
0	Universal Studios Amblin Entertainment Legenda...	6/9/15	5562	
1	Village Roadshow Pictures Kennedy Miller Produ...	5/13/15	6185	
2	Summit Entertainment Mandeville Films Red Wago...	3/18/15	2480	
3	Lucasfilm Truenorth Productions Bad Robot	12/15/15	5292	
4	Universal Pictures Original Film Media Rights ...	4/1/15	2947	

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

To see how many rows and columns it has

```
In [21]: df.shape
```

```
Out[21]: (10866, 21)
```

To look at the data by its count, mean, standard deviation, ...

```
In [22]: df.describe()
```

```
Out[22]:
```

	id	popularity	budget	revenue	runtime
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

Taking a look at the data info to see how many cell are missing, NA or zeros

```
In [23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
```

```

popularity          10866 non-null float64
budget              10866 non-null int64
revenue             10866 non-null int64
original_title      10866 non-null object
cast                10790 non-null object
homepage            2936 non-null object
director            10822 non-null object
tagline             8042 non-null object
keywords            9373 non-null object
overview            10862 non-null object
runtime             10866 non-null int64
genres              10843 non-null object
production_companies 9836 non-null object
release_date        10866 non-null object
vote_count          10866 non-null int64
vote_average        10866 non-null float64
release_year        10866 non-null int64
budget_adj          10866 non-null float64
revenue_adj         10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

### 1.1.2 Data Cleaning (Replace this with more specific notes!)

Dropping some unnecessary columns

```
In [24]: df.drop(['id', 'imdb_id', 'original_title', 'cast', 'homepage', 'director', 'tagline'],
```

To see if these columns are all gone

```
In [25]: df.head()
```

```

Out[25]:   popularity    budget    revenue  runtime  \
0   32.985763  150000000  1513528810     124
1   28.419936  150000000   378436354     120
2   13.112507  110000000   295238201     119
3   11.173104  200000000  2068178225     136
4    9.335014  190000000  1506249360     137

          genres  vote_count  vote_average  \
0  Action|Adventure|Science Fiction|Thriller    5562         6.5
1  Action|Adventure|Science Fiction|Thriller    6185         7.1
2           Adventure|Science Fiction|Thriller    2480         6.3
3  Action|Adventure|Science Fiction|Fantasy    5292         7.5
4           Action|Crime|Thriller    2947         7.3

   release_year
0           2015

```

1	2015
2	2015
3	2015
4	2015

To see how much data is missing

```
In [26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 8 columns):
popularity      10866 non-null float64
budget          10866 non-null int64
revenue         10866 non-null int64
runtime         10866 non-null int64
genres          10843 non-null object
vote_count      10866 non-null int64
vote_average    10866 non-null float64
release_year    10866 non-null int64
dtypes: float64(2), int64(5), object(1)
memory usage: 679.2+ KB
```

Taking a look at the cell where there are NaN values

```
In [27]: df[df.genres.isnull()]
```

```
Out[27]:
```

	popularity	budget	revenue	runtime	genres	vote_count	vote_average	\
424	0.244648	0	0	100	NaN	21	6.1	
620	0.129696	0	0	90	NaN	13	5.0	
997	0.330431	0	0	44	NaN	13	6.8	
1712	0.302095	0	0	88	NaN	57	7.4	
1897	0.020701	0	0	76	NaN	11	7.0	
2370	0.081892	0	0	0	NaN	12	5.8	
2376	0.068411	0	0	62	NaN	11	7.7	
2853	0.130018	0	0	110	NaN	12	7.2	
3279	0.145331	0	0	96	NaN	11	6.1	
4547	0.520520	0	0	220	NaN	12	8.3	
4732	0.235911	0	0	100	NaN	12	6.2	
4797	0.167501	0	0	60	NaN	10	7.8	
4890	0.083202	0	0	2	NaN	14	7.0	
5830	0.248944	0	0	60	NaN	26	8.5	
5934	0.067433	0	0	3	NaN	27	6.9	
6043	0.039080	0	0	127	NaN	12	5.9	
6530	0.092724	0	0	6	NaN	24	5.9	
8234	0.028874	0	0	103	NaN	44	6.7	
8614	0.273934	0	0	12	NaN	14	6.7	
8878	0.038045	0	0	85	NaN	16	5.4	

9307	0.094652	0	0	105	NaN	10	5.3
9799	0.175008	0	0	5	NaN	11	5.0
10659	0.344172	5000	0	71	NaN	10	3.0

	release_year
424	2015
620	2015
997	2014
1712	2009
1897	2009
2370	2010
2376	2010
2853	2001
3279	2008
4547	2012
4732	2012
4797	2012
4890	2012
5830	2013
5934	2013
6043	2013
6530	2005
8234	1995
8614	1996
8878	2000
9307	1989
9799	1974
10659	1970

Removing those NaN cells

```
In [28]: df.dropna(inplace=True)
```

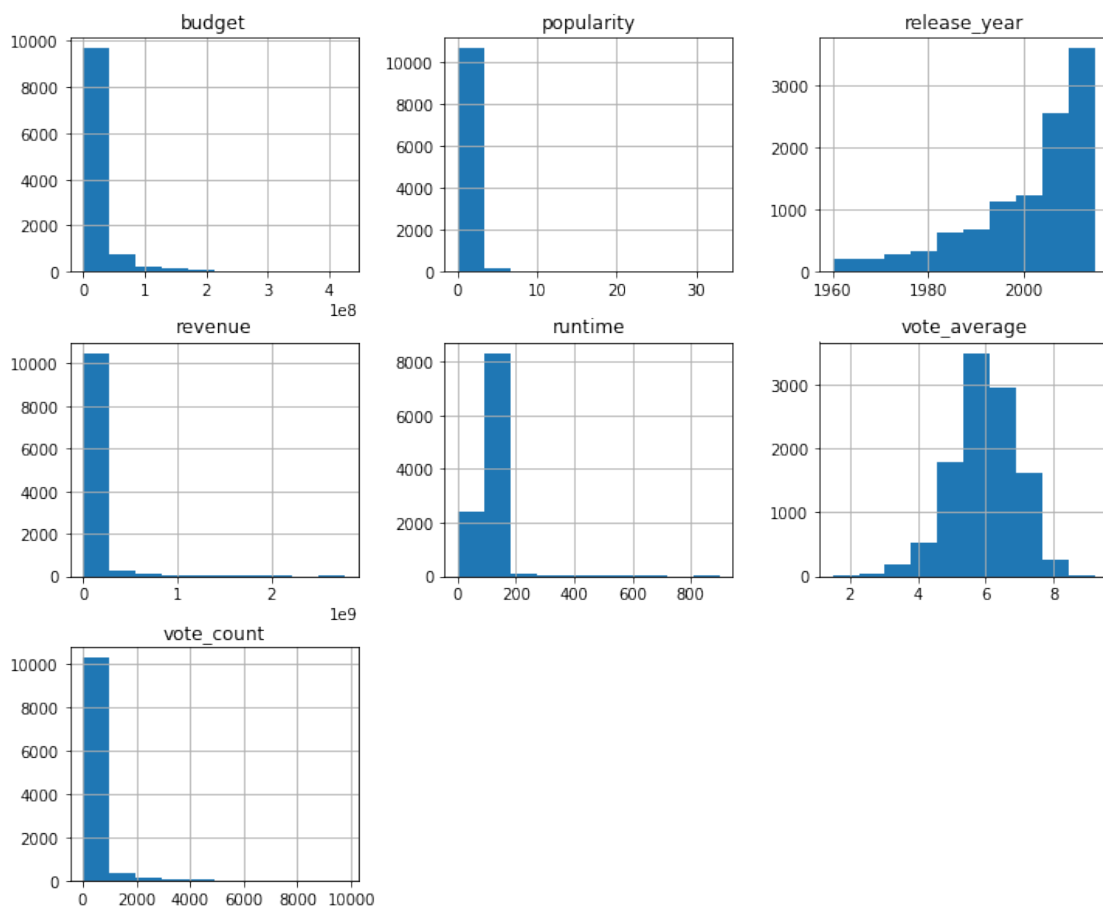
```
In [29]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10843 entries, 0 to 10865
Data columns (total 8 columns):
popularity      10843 non-null float64
budget          10843 non-null int64
revenue         10843 non-null int64
runtime         10843 non-null int64
genres          10843 non-null object
vote_count      10843 non-null int64
vote_average    10843 non-null float64
release_year    10843 non-null int64
dtypes: float64(2), int64(5), object(1)
memory usage: 762.4+ KB
```

Taking a general look at the data by histograms

```
In [30]: df.hist(figsize=(12,10))
```

```
Out[30]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f53a1b6dcc0>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f53a584b198>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f53a5806198>],  
               [[<matplotlib.axes._subplots.AxesSubplot object at 0x7f53a57bd208>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f53a5763c50>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f53a5763c88>],  
               [[<matplotlib.axes._subplots.AxesSubplot object at 0x7f53a56e2160>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f53a56991d0>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x7f53a56d30f0>]], dtype=object)
```




## Exploratory Data Analysis

### 1.1.3 Research Question 1: How did the customer rating of those movies change from the past to present?

I create a scatter plot to show the relation between 'release\_year' and 'vote\_average'. It is better to use scatter plot than any other kinds of graph for this situation

```
In [37]: df1 = df.plot(x='release_year', y = 'vote_average', kind = 'scatter',title='Movie rating
df1.set_ylabel('Vote Average (by star)')
df1.set_xlabel('Releasing Year')
```

```
Out[37]: Text(0.5,0,'Releasing Year')
```



Answer: As illustrated in the scatter graph, recently released movies come with both higher and lower range of rating point compared to the old movies. In other words, customer rating are now more diversity than that of the past. More ever, in the other hand, it seems like we have more worse movie than better movies releasing than that of the past

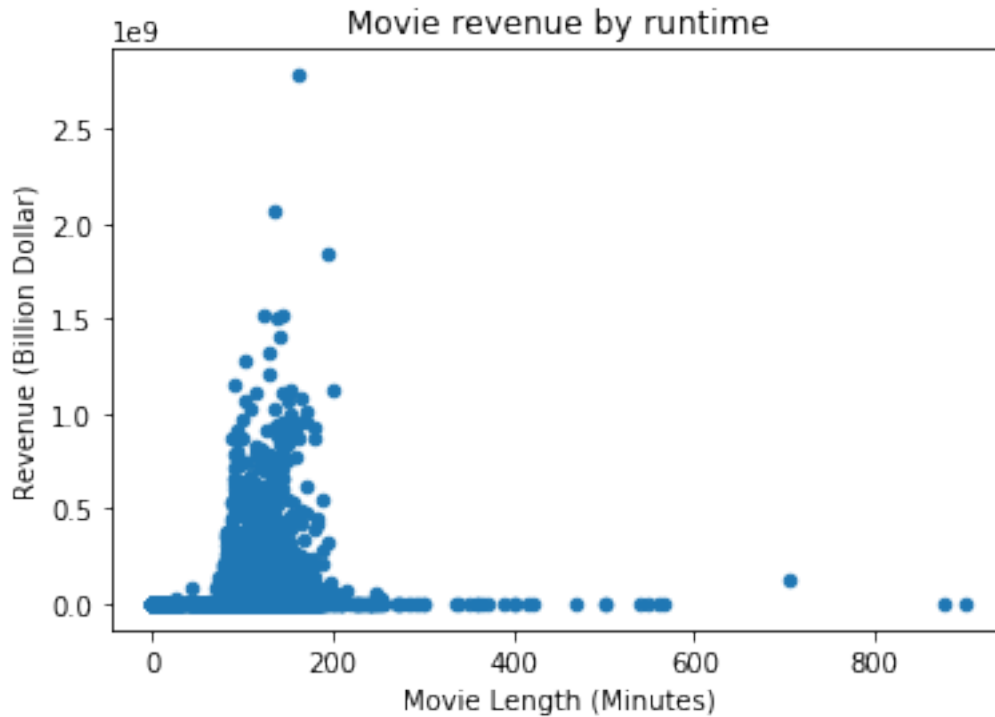
#### 1.1.4 Research Question 2: What is the sweet spot for a movie length? (of which they make the most revenue)

To answer this question, I demostate the relation between 'runtime' and 'revenue' by this graph. It is better to use scatter plot than any other kinds of graph for this situation

```
In [41]: df2 = df.plot(x = 'runtime', y='revenue', kind='scatter', title='Movie revenue by runti
df2.set_xlabel('Movie Length (Minutes)')
df2.set_ylabel('Revenue (Billion Dollar)')
```

```
Out[41]: Text(0,0.5,'Revenue (Billion Dollar)')
```





As demonstrated by the graph, the most profitable movies were created in between 130 to 200 minutes. This is where the dots stay at the highest density level

## Conclusions

The dataset seems like providing me enough information to answer my questions.

Based on my investigation, customer rating are now more diversity than that of the past with the old movie rating average is a little higher than nowadays movie.

To get the most of of profit, a movie should go between 130 to 200 minutes. This charateristic are pretty accurate because it based on a large amount of movies investigated

Convert file command:

```
In [150]: from subprocess import call
          call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
Out[150]: 0
```